

An Exact Fold Condition for Mass-Balanced Models of Protein Quality Control

Kiran Boggavarapu^{1*}

^{1*}Department of Chemistry and Physics, McNeese State University,
Lake Charles, LA, 70609, USA, ORCID 0000-0003-0751-6459.

Corresponding author(s). E-mail(s): kiran@mcneese.edu;

Abstract

A cell clearing misfolded protein with finite pools of chaperones and proteases tolerates damage influx only up to a threshold, beyond which no low-burden steady state exists. It is normally located by continuation, one parameter set at a time. We show it is algebraic. For any clearance network in which the damage influx is state-independent and mass balance accounts for all outflow, the Jacobian determinant factors as $\det J = -(\nabla R \times \nabla G)$, where R is total removal flux and G the aggregate nullcline field. A fold is therefore exactly a constrained critical point of removal on the aggregate nullcline, with critical influx $j_{crit} = R(u^*, a^*)$; the converse holds under three genericity conditions and a fourth that follows from the first. The identity generalises to arbitrary state dimension and survives growth dilution and self-damaging clearance machinery. At the fold the machinery delivers 8.3% of the flux its pools could support, so the capacity ceiling overestimates the tolerable influx about twelvefold under the stipulated catalytic closure and 2.8-fold under the standard bound-complex alternative. Two results bound it: the fold equations return a candidate set with more than one element in 9% of a wide kinetic ensemble, and in 3.9% a Hopf bifurcation precedes the fold, opening a band of influx inside which integration finds sustained oscillation. For aggregation asymmetry in dividing *Escherichia coli*, the framework yields a falsifiable requirement on the old-pole aggregate load, indexed by the share held in the visible polar focus.

Keywords: proteostasis, saddle-node bifurcation, conserved resource, clearance network, growth dilution, Hopf bifurcation

MSC Classification: 92C40 (primary); 37G10 , 34C23 , 92C42 (secondary)

1 Introduction

Protein quality control is a resource-limited clearance problem. Chaperones refold damaged monomers, proteases degrade them, and both draw on pools that are finite and shared across substrates (Hipp et al. 2019). Damage arrives continuously from translation error, thermal stress and oxidative modification. When influx exceeds what the machinery can dispose of, aggregate accumulates and the cell loses viability.

Models built on this premise locate the threshold numerically. One fixes a parameter set, integrates, continues the stable branch in the influx and records where it terminates. Cotton et al. (2026) illustrate both the power and the cost of this approach: across eight variants of aggregate formation and removal kinetics they demonstrate that the same bifurcation structure is preserved, with a stable low-aggregate state and an unstable upper branch merging at a critical monomer concentration. The universality is established by exhibiting it case by case. Nothing in that treatment says why it should hold, or what would have to change for it to fail.

We show that for one structural class the threshold is algebraic rather than computational. Two facts about clearance models — the damage influx enters one state equation, and mass balance is exact — force the Jacobian determinant to factor into the cross product of two gradients. A fold is then a constrained critical point of total removal, in the Lagrange sense, on the nullcline of the non-influx equation, and the critical influx is the value of removal at that point. Adding states, adding regulation, adding cell division, or making the clearance capacity itself depend on the load leaves the identity intact.

The hypothesis required is weaker than it first appears. The theorem does not need influx and clearance capacity to be independent. It needs only that total influx be state-independent and that mass balance count all outflow. Capacity may depend on the error rate, on the burden, or on both. This matters because in a cell the chaperones and proteases are themselves translated by the error-prone apparatus whose product they clear.

Two results bound what the fold condition delivers. Solving the fold equations returns the candidate set exactly and without a sweep, but in about 9% of a wide kinetic ensemble that set has more than one element, and identifying which candidate terminates the accessible branch requires information the equations do not carry. And the fold is not always the first loss of stability: in the same ensemble a Hopf bifurcation precedes it in a characterisable region of parameter space.

Section 2 fixes the model class and its hypotheses. Section 3 states and proves the fold condition, gives the genericity conditions for the converse, places the result relative to the homeostasis literature, and treats orientation and multiplicity. Section 4 gives three extensions. Section 5 reports numerical verification. Section 6 characterises where the fold sits relative to the capacity ceiling. Section 7 treats instabilities that precede the fold. Section 8 applies the framework to aggregation asymmetry in dividing *E. coli* and derives the requirement on old-pole aggregate load that is the paper’s one directly falsifiable prediction. Section 9 states what the theory predicts and what would refute it, and Section 10 says what the result does and does not deliver.

2 The model class

2.1 State equations

Let u denote damaged or misfolded soluble monomer and a aggregate burden. Both are **total** pools — free plus machinery-bound — carried in monomer-equivalent units, so that transfer between them conserves mass one for one. The baseline field is

$$\begin{aligned}\frac{du}{dt} &= j - v_{ref} - v_{degU} - n - g + v_{dis}, \\ \frac{da}{dt} &= n + g - v_{dis} - v_{degA}.\end{aligned}$$

Chaperone and protease pools are conserved and shared, and every rate law is written in free concentrations. With 1:1 rapid-equilibrium complexes CU , CA , CN , DU and DA , the four conservation laws close simultaneously on the free concentrations:

$$\begin{aligned}u_f &= \frac{u}{1 + C_f/K_{CU} + D_f/K_{DU}}, & C_f &= \frac{C_{tot}}{1 + \nu + u_f/K_{CU} + a_f/K_{CA}}, \\ a_f &= \frac{a}{1 + C_f/K_{CA} + D_f/K_{DA}}, & D_f &= \frac{D_{tot}}{1 + u_f/K_{DU} + a_f/K_{DA}}.\end{aligned}$$

Given rapid equilibrium these four equations are identities rather than approximations, and they are solved jointly at every state; total substrate is never substituted into a formula expecting a free concentration. Here $\nu = N_f/K_N$ is the occupancy contributed by nascent chains, which consume capacity without contributing damage influx — the mechanism by which ordinary translational load can move the fold.

R and G are therefore defined implicitly rather than by formulae in (u, a) , and their regularity has to be established rather than assumed.

Lemma 0 (well-posedness of the closure). *For every (u, a) in the nonnegative quadrant the four equations above have exactly one solution, and it is strictly positive. The binding Jacobian in (C_f, D_f) is a nonsingular M -matrix with $\det \geq 1 + \nu$, so the free concentrations are real-analytic in (u, a) on the open quadrant. Consequently $\partial u_f / \partial u > 0$ and $\partial a_f / \partial u \geq 0$, while both free machinery pools are nonincreasing in u .*

Proof. The map sending (C_f, D_f) to the right-hand sides of the two pool equations carries the box $[0, C_{tot}] \times [0, D_{tot}]$ into itself and is monotone increasing, since raising either free pool lowers both free substrates and so lowers both denominators; Knaster–Tarski gives a least and a greatest fixed point. It is also strictly sublinear — scaling the argument by $\lambda < 1$ shrinks the occupancy terms of each denominator by λ but leaves the additive constant, so the quotient rises by more than λ — and a monotone strictly sublinear map has at most one positive fixed point, by the standard argument on the largest t with $x \geq ty$. For the Jacobian, all four partials of the free substrates in the free pools are negative, so the off-diagonal entries are nonpositive. Contracting against the positive vector (C_f, D_f) and using that each free substrate is homogeneous of degree -1 in $(1, C_f, D_f)$ gives row values $C_f(1 + \nu + s_u^{-1}u_f/K_{CU} + s_a^{-1}a_f/K_{CA})$

and $D_f(1 + s_u^{-1}u_f/K_{DU} + s_a^{-1}a_f/K_{DA})$, both positive, where s_u and s_a are the two substrate denominators. A matrix with nonpositive off-diagonal entries admitting a positive vector of positive image is a nonsingular M-matrix, whence $\det \geq 1 + \nu$ and the inverse is nonnegative. The implicit function theorem then applies at every state, and inverse-positivity applied to the derivative of the closure in u gives the two monotonicity statements. $\square$

The rate laws are

$$\begin{aligned}
v_{ref} &= \frac{k_{ref} C_f u_f}{K_{ref} + u_f} && \text{refolding} \\
v_{degU} &= \frac{k_U D_f u_f}{K_U + u_f} && \text{soluble degradation} \\
n &= k_n u_f^m, \quad m > 1 && \text{nucleation} \\
g &= k_g u_f a_f && \text{aggregate growth} \\
v_{dis} &= \frac{k_{dis} C_f a_f}{K_{dis} + a_f} && \text{disaggregation} \\
v_{degA} &= \frac{k_A D_f a_f}{K_A + a_f} && \text{aggregate clearance}
\end{aligned}$$

Time is measured in units of $1/k_{ref}$, so refolding carries unit coefficient and the remaining constants enter as the ratios $\rho_U = k_U/k_{ref}$, $\rho_A = k_A/k_{ref}$, $\alpha_n = k_n/k_{ref}$, $\alpha_g = k_g/k_{ref}$ and $\alpha_d = k_{dis}/k_{ref}$, with κ_{ref} , κ_u , κ_a and κ_{dis} for K_{ref} , K_U , K_A and K_{dis} . Sections 3.3, 6 and 7 quote the scaled names.

The two aggregate-forming terms are distinct processes. $n = k_n u_f^m$ is primary nucleation with an effective reaction order greater than one, in the sense in which fitted orders are understood as coarse-grained descriptions of multi-step processes (Ferrone 1999; Knowles et al. 2009; Meisl et al. 2017). $g = k_g u_f a_f$ is elongation: growth of existing aggregate at a monomer-dependent rate. The model carries no secondary nucleation term, which would be autocatalytic in aggregate surface (Cohen et al. 2013); adding one changes the coefficients below but not the structure, since it enters G and not R .

The dissociation constants K_{CU} and K_{CA} describe competitive occupancy of the pool; the Michaelis constants K_{ref} , K_U , K_{dis} and K_A describe a rate-limiting catalytic step acting on the pool not held in any complex. Treating them as successive steps is a stipulation and not a derivation. Written against the bound complex instead, each flux would carry no second saturation, so the coded form asserts nonproductive sequestration followed by independent productive capture, and gives cycles sharing one pool independent Michaelis factors rather than one competitive denominator. Both stipulations bind. Over the kinetic ensemble of Section 5 the two catalytic occupancies drawing on the chaperone pool sum above unity at 334 of 2767 fold states, and the substrate that productive complexes would hold reaches a fifth of the total soluble pool at the median. Section 4.4 reports what moves under the alternative closure. Nothing in Theorem 1 depends on the choice: the hypotheses below hold for either.

Write the total removal flux and the aggregate field as

$$R(u, a) = v_{ref} + v_{degU} + v_{degA}, \quad G(u, a) = \frac{da}{dt}$$

R is what leaves the damaged compartment; G is the non-influx equation. Note that n and g appear in R nowhere: they transfer material between u and a and cancel from $du/dt + da/dt$. Aggregation does not oppose removal directly. It reshapes the nullcline, redistributes burden between the two pools, and suppresses effective removal through saturation and through sequestration of the shared machinery. Section 6 separates those two contributions.

2.2 Hypotheses

H1 (state-independent influx). The damage influx j is a parameter appearing additively in du/dt . Two consequences are used separately and are worth separating. **(H1a)** Total influx does not depend on the state, which is what puts mass balance in the form H2 requires. **(H1b)** j does not appear in da/dt , so the curve $\{G = 0\}$ does not move with the load. Only (H1a) enters the identity. Section 4.3 relaxes (H1b) by letting clearance capacity depend on the influx, and Theorem 1.4 covers that case.

H2 (exact mass balance). Internal transfer between u and a cancels between the two equations, so $du/dt + da/dt = j - R$ holds identically.

H3 (smoothness). R and G are continuously differentiable on the region of interest. For the model class of Section 2.1 this is a consequence rather than an assumption: Lemma 0 makes the free concentrations real-analytic in the state, and every rate law is rational in them apart from $k_n u_f^m$, which is C^1 at the origin for $m > 1$. H3 is retained as a hypothesis because Theorem 1 is stated for the wider class in which the fields are given directly.

H4 (deterministic, well-mixed, finite-dimensional). The system is an ODE in concentrations.

Remark 1. H1 constrains the influx, not the capacity. Nothing forbids C_{tot} or D_{tot} from depending on j or on the state. The gradients in Section 3 are taken at fixed j , so parameter dependence on the influx cannot enter them; the row operation needs (H1a) and H2 alone, and neither mentions da/dt . What capacity feedback costs is (H1b), and with it the fixed nullcline that makes the candidate solve cheap. Corollary 4 treats the consequences of capacity that degrades with load.

Remark 2. H4 is false in the system of Section 8 in two respects. Aggregates in *E. coli* are spatially segregated into polar inclusion bodies rather than well mixed, and molecule numbers per damaged species are small. Both effects act in the same direction: noise-induced escape precedes the deterministic fold, so a measured threshold should sit below j_{crit} and be smeared across a population rather than sharp.

3 The fold condition

3.1 The identity

Four statements are separated because they need different hypotheses; they are referred to collectively as Theorem 1.

Theorem 1.1 (mass-balance identity). *Assume (H1a), H2 and H3. Then at fixed j*

$$\det J = -(\nabla R \times \nabla G)$$

identically, the gradients being taken with respect to the state.

Proof. By H2, $du/dt + da/dt = j - R$ exactly. The determinant-preserving row operation $\text{row}_1 \rightarrow \text{row}_1 + \text{row}_2$ replaces the first row by the gradient of $j - R$, and j is a parameter, so

$$\begin{aligned} \det J &= \det \begin{bmatrix} \partial(du/dt)/\partial u & \partial(du/dt)/\partial a \\ \partial(da/dt)/\partial u & \partial(da/dt)/\partial a \end{bmatrix} = \det \begin{bmatrix} -R_u & -R_a \\ G_u & G_a \end{bmatrix} \\ &= -(R_u G_a - R_a G_u) = -(\nabla R \times \nabla G). \quad \square \end{aligned}$$

(H1b) is not used, and the identity is indifferent to whether j enters G or the capacity depends on the load, because the gradients are taken at fixed influx.

Theorem 1.2 (removal along the nullcline). *Assume in addition that $\nabla G \neq 0$. Parametrise $\{G = 0\}$ by arclength s with unit tangent $\gamma' = (-G_a, G_u)/\|\nabla G\|$, and put $r(s) = R(\gamma(s))$. Then*

$$r'(s) = \nabla R \cdot \gamma' = \frac{\det J}{\|\nabla G\|}$$

at every point of the nullcline, not only where it vanishes.

$\det J$ is the derivative of total removal along the aggregate nullcline, up to the positive factor $\|\nabla G\|$. The informal reading — the fold is where total removal stops responding to burden along the nullcline — is therefore an exact pointwise identity rather than a gloss on a vanishing condition, and $\det J = 0$ is $r'(s) = 0$, a constrained critical point of R on $\{G = 0\}$ with critical value $j_{crit} = R(u^*, a^*)$. Both statements are at fixed influx, which is all (H1a) supplies; the curve and the function r may still move with j .

Regularity of the constraint is a hypothesis of the Lagrange reading and not only of the converse. Where $\nabla G = 0$ the determinant vanishes automatically, $\{G = 0\}$ need not be a manifold, and there is no tangent to differentiate along; the two networks discussed in Section 3.2 are that case, not counterexamples to anything.

Theorem 1.3 (fixed-nullcline shortcut). *Assume (H1b), so $\partial G/\partial j = 0$. Then $\{G = 0\}$ is independent of the influx — changing the load slides the state along it — equilibria are exactly its points with $R = j$, and the fold candidates solve $\{G = 0, \det J = 0\}$ in the state alone.*

This is the whole practical content: neither equation contains j , so locating candidates is a two-dimensional root solve rather than a continuation sweep. It is also what makes the branch one-dimensional in the useful sense — with the curve and r both fixed, the equilibria on the nullcline are the solutions of $r(s) = j$, so the equilibrium branch is the level-set inverse of one scalar function of one variable, and its turning points in j are the critical points of r .

Theorem 1.4 (load-dependent capacity). *When $\partial G/\partial j \neq 0$ the nullcline moves with the influx and $j_{crit} = R(u^*, a^*)$ becomes a self-consistency condition; candidates solve $\{G = 0, \det J = 0, R = j\}$ in (u, a, j) . Theorem 1.1 is unaffected.*

At the base parameter set the alignment is visible rather than inferred: the removal contours meet the aggregate nullcline tangentially at the solved fold, to a sine of 3.54×10^{-10} (Fig. 1a). The same computed branch separates two singularities that are easily conflated. The soluble coordinate has a horizontal tangent at $j_{turn} = 0.154090$, where $G_a = 0$ and the nullcline runs vertical in the burden plane; that point is regular, since $\det J = R_a \cdot G_u = 2.027 \times 10^{-3}$ there. The fold is at $j_{crit} = 0.154239$, where both coordinates turn together (Fig. 1b). Here they differ by one part in a thousand, with j_{turn} the smaller. This is also why tracing $\{G = 0\}$ by solving for a at fixed u fails: the two roots in a merge exactly at j_{turn} , returning the curve in two disconnected pieces with the fold in the gap between them. The root-finder's failure and the horizontal tangent are one locus, not two problems.

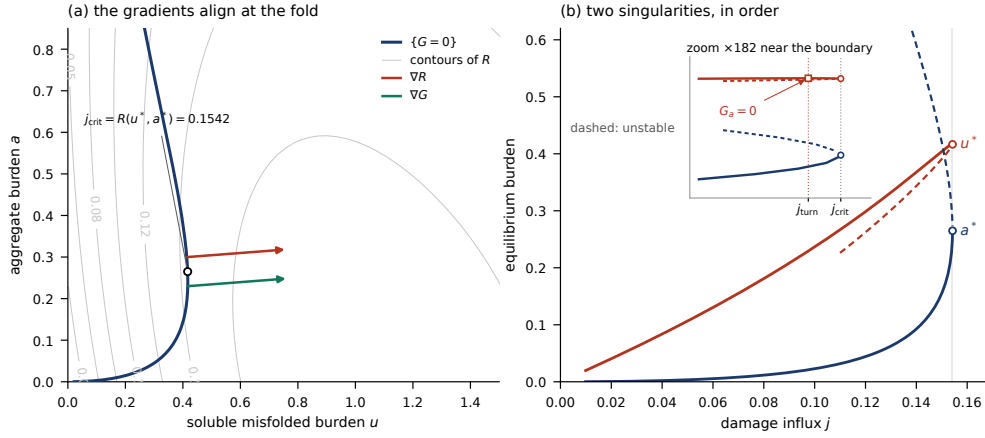


Fig. 1 The fold as a constrained critical point. **(a)** Phase plane at the base parameter set. The aggregate nullcline $\{G = 0\}$ (dark) is fixed under change of influx; grey contours are total removal R . At the solved fold $(u^*, a^*) = (0.4166, 0.2650)$ the gradients ∇R and ∇G are parallel, which is the Lagrange condition of Theorem 1.2; they are drawn from offset origins because a common origin would hide one behind the other. The sine of the angle between them is 3.54×10^{-10} , and $j_{crit} = R(u^*, a^*) = 0.1542$. **(b)** Both equilibrium coordinates against influx along the same curve. Cramer's rule gives $du^*/dj = -G_a/\det J$ and $da^*/dj = G_u/\det J$, so the branch carries two distinct, ordered singularities. At $j_{turn} = 0.154090$ the soluble coordinate has a horizontal tangent where $G_a = 0$; this is a regular point, $\det J = R_a \cdot G_u = 2.027 \times 10^{-3}$ there rather than zero, and it lies on the stable branch at $j_{turn}/j_{crit} = 0.9990$. At $j_{crit} = 0.154239$ both coordinates have a vertical tangent, which is $\det J = 0$. The inset is a $\times 182$ zoom on the same computed data, not a schematic. The aggregate coordinate has no horizontal tangent anywhere: $G_u > 0$ at all 305 traced points, minimum 2.077×10^{-3} .

3.2 Genericity, and the converse

Theorem 1.2 gives one direction: a fold satisfies gradient alignment. The converse requires that the degeneracy be a fold rather than a transcritical or pitchfork bifurcation, a cusp, a double-zero eigenvalue, or a curve of equilibria. Three conditions suffice:

- **(G1)** regularity of the constraint, $\nabla G \neq 0$;
- **(G2)** a simple zero eigenvalue, $\text{tr } J \neq 0$;
- **(G3)** nondegeneracy, $d^2 R/ds^2 \neq 0$ along the nullcline at the critical point.

Parameter transversality, $w \cdot \partial F/\partial j \neq 0$ for the left null vector w , is the fourth condition of the classical statement. It is not independent here, and the corollary below shows it follows from (H1b) and (G1).

Theorem 2 (converse). *Let F be C^2 and assume (H1b). Under (G1)–(G3), a constrained critical point of R on $\{G = 0\}$ is a generic saddle-node in j , with $j_{crit} = R(u^*, a^*)$.*

Proof. By (G1) the nullcline is a C^2 curve near the point, so Theorem 1.2 applies and $\det J = 0$ is $r'(0) = 0$. With (G3),

$$r(s) = j_{crit} + \frac{1}{2}r''(0)s^2 + O(s^3),$$

a nondegenerate extremum, so for j on one side of j_{crit} there are exactly two equilibria near the point and on the other side none: the pair collides and annihilates. Three conditions then remain. By (G2) the zero eigenvalue is algebraically simple, since in two dimensions the other eigenvalue is $\text{tr } J$, which also places it off the imaginary axis. By (H1b) and (G1) the parameter transversality $w \cdot \partial F/\partial j \neq 0$ holds, by the corollary below. For the quadratic nondegeneracy, let v span the kernel of J ; the second row of $Jv = 0$ reads $\nabla G \cdot v = 0$, so v is tangent to $\{G = 0\}$ and parallel to $\gamma'(0)$. With $w = (1, 1 + \lambda)$ the left null vector and $\nabla R = \lambda \nabla G$,

$$w \cdot D^2 F(v, v) = -D^2 R(v, v) + \lambda D^2 G(v, v),$$

while differentiating $G(\gamma(s)) = 0$ twice gives $\nabla G \cdot \gamma'' = -D^2 G(v, v)$, so that $r''(0) = D^2 R(v, v) - \lambda D^2 G(v, v)$. Hence

$$r''(0) = -w \cdot D^2 F(v, v),$$

and (G3) is not merely equivalent to the classical nondegeneracy condition of the saddle-node but equal to it up to sign. The three hypotheses of Sotomayor's theorem are met, and it supplies the normal form, with quadratic coefficient of sign $-r''(0)$. $\square$

The nullcline is not a centre manifold and the proof does not use it as one. $\{G = 0\}$ is not invariant: along the flow $dG/dt = G_u \cdot (j - R)$, which vanishes only at the equilibrium, so trajectories cross the curve rather than run along it. What the tangency supplies is that the arclength direction and the centre direction agree at the critical point, which is why $r''(0)$ and the centre-manifold quadratic coefficient are the same number. In the suspended system with $dj/dt = 0$ the curve of equilibria does lie in the extended centre manifold; the fixed-influx nullcline does not.

Two things come free. The sign of $r''(0)$ is the orientation of Section 3.4, so that classification is a corollary rather than a separate observation. And nothing in the

proof appeals to numerical verification; Table 1 reports whether the hypotheses hold in the ensembles studied, which is a question about the model rather than about the theorem.

Remark 3 (what a fold terminates). Step two of the proof also classifies the collision. Since the eigenvalues at a fold are 0 and $\text{tr } J$, a negative trace means a stable node and a saddle merge, so the fold terminates a stable branch and is the collapse threshold. A positive trace means an unstable node and a saddle merge, and the low-burden state has already lost stability somewhere below: the fold is a turning point of the branch but not a boundary of viability. On the load grid every fold has $\text{tr } J < 0$. Over the kinetic box 61 of 2767 have $\text{tr } J > 0$, and all 61 are among the networks of Section 7 whose branch loses stability before the fold — as they must be, which makes the count a consistency check between two independent computations rather than a finding. The finding is the converse: of the 108 networks that lose stability first, 47 recover it before reaching the fold — 38 of them with both crossings resolved — so an oscillatory excursion below j_{crit} does not imply that the fold itself is unstable. Section 7 treats the two groups separately.

Corollary (transversality is automatic). *Under (H1b) and (G1), $w \cdot \partial F / \partial j \neq 0$.*

Proof. By (H1b), $\partial F / \partial j = (1, 0)$ exactly, so the condition is $w_1 \neq 0$. Let $L = I + e_1 e_2^T$, so that LJ is the matrix whose first row is $-\nabla R$ and whose remaining rows are the non-influx constraint gradients, and note that L fixes the first coordinate of any row vector multiplying it on the left. Put $z = wL^{-1}$, so $z(LJ) = 0$. If $z_1 = 0$ the remaining components annihilate the non-influx constraint gradients, which have full rank by (G1), forcing $z = 0$ and hence $w = 0$. So $z_1 \neq 0$, and $w_1 = z_1 \neq 0$. $\square$

The argument uses only $\partial F / \partial j = e_1$ and full rank, so it holds in every state dimension. It also gives the margin in closed form: normalising $w = (1, 1 + \lambda)$ with $\nabla R = \lambda \nabla G$,

$$\frac{|w_1|}{\|w\|} = (1 + (1 + \lambda)^2)^{-1/2},$$

which is small exactly when $|\lambda|$ is large — that is, when $\nabla G \rightarrow 0$ at fixed ∇R . The near-failures of the constraint and of transversality are therefore one event and not two, which is what Table 1’s two exceptions are.

Computing the two sides of the closed form by independent routes — the left from a singular value decomposition of the analytic Jacobian, the right from central-difference gradients — they agree at all 325 load-grid states and at 2765 of the 2767 kinetic-box states, at median relative difference 8.5×10^{-11} and maximum 1.8×10^{-5} . The two that disagree are precisely the two exceptions of Table 1, where $\|\nabla G\|$ is 1.04×10^{-8} and 3.99×10^{-8} against a smallest value of 1.41×10^{-6} everywhere else: there the gradient ratio is differencing noise and the closed form has nothing to evaluate.

Table 1 The genericity conditions of Theorem 2, evaluated at every solved fold state in each ensemble. The first three rows are the hypotheses; each entry is the minimum of that condition’s margin over the whole population, so a value bounded away from zero means the condition holds throughout. The fourth row is a consequence of the first rather than an independent condition, and takes the closed form given above wherever the constraint is regular; its two exceptions are the same two networks

as (G1)’s, which is forced. The last row reports a sign rather than a margin, because $|\text{tr } J|$ conceals the classification: the eigenvalues at a fold are 0 and $\text{tr } J$, so a negative trace is a stable node colliding with a saddle and the fold terminates a stable branch, while a positive trace is an unstable node colliding with a saddle and the fold is not a stability boundary at all.

condition	load grid (325)	kinetic box (2767)
(G1) $ \nabla G $	min 0.106	min 1.04×10^{-8} , 2 exceptions
(G2) $ \text{tr } J $	min 0.303	min 8.38×10^{-4}
(G3) $ d^2 R/ds^2 $	min 0.0929	min 9.21×10^{-6}
consequence of (G1), $ w_1 / w $	min 0.341	min 3.34×10^{-9} , 2 exceptions
sign of $\text{tr } J$	all negative, -0.365 to -0.303	-626 to $+3.41$; positive at 61

The two exceptions are the same two networks, and they fail (G1) and (G3) for one reason, taking the transversality margin down with them: both sit at aggregate burden of order 10^{-11} and 10^{-15} , where the aggregate compartment is empty, $\nabla G \rightarrow 0$, the constraint ceases to be a manifold, and every condition defined on it degenerates. At both, $\text{tr } J = -\nabla R$ exactly, the signature of a state where G has gone flat. These are the boundary of the model’s domain of definition rather than counterexamples to Theorem 2. A further 117 of 2884 sampled networks carry no solvable state satisfying $\{G = 0, \det J = 0\}$ and are excluded from all conditions above.

Both computations use arclength along the nullcline rather than the u -coordinate. The nullcline turns vertical where $G_a = 0$, which is a coordinate singularity of the u -parametrisation, and every quantity evaluated on the curve is affected there.

3.3 Relation to infinitesimal homeostasis

The hypotheses of Theorem 1 nearly coincide with those of an existing programme, and the relation is worth stating precisely.

Wang et al. (2021) study input-output networks in which a parameter enters exactly one node’s equation — the condition H1 restates — and ask when the input-output function $x_o(I)$ has vanishing derivative. That question has a long history in the adaptation literature, where the stronger requirement of perfect adaptation demands the derivative vanish identically over an interval and the capable network topologies have been characterised numerically and then structurally (Ma et al. 2009; Ferrell 2016; Araujo and Liotta 2018; Khammash 2021). Cramer’s rule applied to the implicitly differentiated equilibrium condition gives $x'_o = \pm f_{i,I} \det(H)/\det(J)$, where the homeostasis matrix H is the Jacobian with its first row and last column deleted. Infinitesimal homeostasis is $\det(H) = 0$, and the programme built on that identity applies the Frobenius-König theorem to factor the determinant into irreducible blocks, each associated with a subnetwork motif of structural or appendage type (Golubitsky and Stewart 2017; Reed et al. 2017; Golubitsky and Wang 2020; Huang and Golubitsky 2022; Madeira and Antoneli 2024; Antoneli et al. 2025; Lin et al. 2026). The framework has since been extended to mass-action systems obeying conservation laws, using

reduced Jacobian and reduced homeostasis matrices that incorporate the conserved quantities (Jin and Rempala 2026).

The two conditions are complementary degeneracies of the same expression. Theirs is the vanishing of the numerator; ours of the denominator. Their derivation presupposes $\det J \neq 0$, since the equilibrium is assumed hyperbolic, so the case Theorem 1 characterises is excluded from their setting by hypothesis. Geometrically, infinitesimal homeostasis is a horizontal tangent of the equilibrium branch in the input-output plane, where the output stops responding to the input; a fold is a vertical tangent, where the branch turns and the implicit function theorem fails. Neither condition implies the other.

Three degeneracies of the same equilibrium branch therefore arise: $\det H = 0$ is homeostasis, $\det J = 0$ with $\text{tr } J \neq 0$ is the fold, and $\text{tr } J = 0$ with $\det J > 0$ is a Hopf bifurcation. This paper exhibits computed instances of the second and third (Sections 5 and 7) and a structural exclusion of the first. Taking the aggregate as output, the network is two nodes with a single arrow, H is 1×1 , $\det H = G_u$, and the only homeostasis type available is the degree-one Haldane form.

It does not occur, and the reason reduces to the sign of one term. Differentiating through the closure, G_u is a sum of four products of which three are nonnegative at every admissible state: nucleation and elongation both rise with u , while raising u titrates chaperone and protease away from disaggregation and aggregate clearance, which enter G negatively and whose free pools are nonincreasing in u by the inverse-positivity of Lemma 0. The fourth carries the sign of $\partial G / \partial a_f$,

$$G_{af} = \alpha_g u_f - \alpha_d c_f \frac{\kappa_{dis}}{(\kappa_{dis} + a_f)^2} - \rho_A d_f \frac{\kappa_a}{(\kappa_a + a_f)^2},$$

so $G_{af} \geq 0$ is sufficient for $G_u > 0$.

That inequality is Section 7's oscillatory region written as a condition rather than as a description. Clearance that saturates sharply is a small κ_a , which shrinks the third term; aggregation that is growth-dominated rather than nucleation-dominated is a large $\alpha_g u_f$, which is the first. Measured at the 2767 fold states, the two negative terms carry medians of 0.5% and 6.7% of the positive one over the 47 window networks, against 76% and 71% over the 2659 that never cross — the disaggregation term suppressed 139-fold and the clearance term 10.6-fold. $G_{af} \geq 0$ accordingly holds at 93.6% of window networks and 86.9% of terminal ones, against 25.3% of the rest. Haldane homeostasis is excluded provably in exactly the region where oscillation occurs, because both statements say that aggregate is self-amplifying at fixed free machinery.

Where the sufficient condition fails the conclusion still holds, and the residual runs the right way. No state of either ensemble has $G_u \leq 0$. The smallest value anywhere, 3.4×10^{-11} , lies where $G_{af} \geq 0$ and the proof applies, at a state with $a = 1.2 \times 10^{-11}$; among the 1998 states where $G_{af} < 0$ and only the competition of terms remains, the smallest is 7.6×10^{-7} , larger by a factor of 2.3×10^4 .

What makes our determinant factor differently is H2, which has no counterpart in the input-output setting. Mass balance supplies the row operation that replaces the influx row by $-\nabla R$, so $\det J$ factors into gradients of identified physical fluxes rather than into combinatorial blocks. Whether the combinatorial machinery transfers to $\det J = -\det[\nabla R; \nabla G; \nabla C]$ is open; a Frobenius-König factorisation (Brualdi and

Ryser 1991) of that determinant would classify fold mechanisms by network topology in the way structural and appendage blocks classify homeostasis mechanisms, and the decomposition of Section 6 is a two-term, physically motivated version of the same question.

The complementary question — when regulation fails rather than when it holds — was raised early in that literature and developed less. Nijhout et al. (2014) ask what happens when a homeostatic mechanism is pushed past its range, and Reed et al. (2018) treat homeostasis at equilibria that are not stable. Neither is a singularity-theoretic account of the failure. Theorem 1 supplies one for the case in which the failure is a fold.

The clearance-limited aggregation literature approaches the same object from the modelling side. Thompson et al. (2021) analyse the stability of a minimal aggregate-removal model; Brennan et al. (2023) treat clearance in a spatial setting; Cotton et al. (2026) generalise to a phase-plane framework across eight aggregation and removal mechanisms, extracting two sufficient conditions for bistability and seeding — self-replicating aggregation with a monomer-dependent rate, and limited-capacity removal. Their input parameter is the monomer concentration, which enters every term of both moment equations, so H1 fails for their system and Theorem 1 does not apply to it as stated. The model classes are adjacent rather than nested. What the present result adds within its own class is a derivation of why the bifurcation structure is preserved under extension, rather than a demonstration that it is.

3.4 Orientation and multiplicity

Two properties of the candidate set follow from the theorem and bound what it delivers.

Orientation. The sign of d^2R/ds^2 distinguishes two kinds of fold. Where it is negative, R has a constrained local maximum, $R = j$ has solutions below $R(x^*)$ and none above, and a pair of equilibria is destroyed as the influx rises: a collapse-oriented fold. Where it is positive, the pair is created as the influx rises: a birth-oriented fold, which is the lower turning point of a hysteresis loop rather than a collapse point. On the load grid $d^2R/ds^2 < 0$ at all 325 folds. Over the kinetic box, 26 of 2765 are birth-oriented, and all 26 carry a collapse-oriented candidate at strictly higher influx. Against them, a control block of 153 collapse-oriented networks — every eighteenth of the 2739, taken systematically rather than at random — carries the same pair in 7 cases. So the pairing is not peculiar to the 26, and 4.6% is an estimate of its rate among ordinary folds rather than a count of them. Hysteresis is therefore a feature of the model class rather than of the 26. The operative consequence is that solving $\det J = 0$ returns candidates without distinguishing them, so orientation must be computed and not inferred from a solver having converged.

Multiplicity. Theorem 1.3 returns candidates, not a candidate.

Corollary 1. *Solving $\{G = 0, \det J = 0\}$ returns the candidate set exactly and without a sweep. Identifying which candidate terminates the accessible low-burden branch requires branch information in 9.1% of the kinetic ensemble (252 of 2765 networks, up to five candidates).*

4 Extensions

4.1 Arbitrary state dimension

Corollary 2. *Let the state be $x = (u, a, c, \dots)$ with du/dt the only equation containing j . Suppose there is a constant mass vector m with $m_1 = 1$ such that*

$$m^\top F(x, j) = j - R(x, j) \quad \text{identically.}$$

Then, writing the non-influx equations as $F_2, \dots, F_n$,

$$\det J = -\det [\nabla R; \nabla F_2; \dots; \nabla F_n]$$

and, provided the non-influx gradients have full rank, this vanishes exactly when ∇R lies in their span — the Lagrange condition for a constrained critical point of R on the intersection of the non-influx nullclines.

The proof is the row operation $\text{row}_1 \rightarrow \text{row}_1 + \sum_{k \geq 2} m_k \text{row}_k$, which preserves the determinant **because** $m_1 = 1$, and replaces the first row by $\nabla(m^\top F) = \nabla(j - R) = -\nabla R$ at fixed j . The sign is -1 in every dimension. Full rank is a hypothesis of the equivalence and not a refinement of it: where the non-influx gradients are dependent the determinant vanishes for every ∇R and the Lagrange reading says nothing, exactly as $\nabla G = 0$ empties it in the plane.

The mass vector is what makes the corollary cover its own instances. A state that carries damaged material enters with $m_k = 1$; a controller state that carries none enters with $m_k = 0$. The two systems below are $m = (1, 1, 0)$ and $m = (1, 1, 1)$, and the earlier two-term statement $du/dt + da/dt = j - R$ covered only the first. The practical consequence is unchanged: a model in this class can be extended without re-deriving its fold condition.

The converse needs restating above two dimensions, because (G2) does two jobs at once in the plane and neither survives. Write the characteristic polynomial as $p(\lambda) = \lambda^n + \dots + c_1 \lambda + c_0$, so that $c_0 = \pm \det J$ and c_1 is, up to sign, the sum of the $(n-1) \times (n-1)$ principal minors. A zero eigenvalue is $c_0 = 0$; it is algebraically simple exactly when $c_1 \neq 0$; and partial hyperbolicity is the separate requirement that no other eigenvalue lie on the imaginary axis. For $n = 2$ both reduce to $\text{tr } J \neq 0$, since $c_1 = -\text{tr } J$ and the remaining eigenvalue is $\text{tr } J$ — which is why the distinction is invisible in the plane and why (G2) may be stated as the trace condition there. In general the converse requires (G1) with the non-influx constraint gradients of full rank, $c_1 \neq 0$, no other eigenvalue on the imaginary axis, and (G3).

We verify Corollary 2 on two three-state systems. In the first the chaperone pool is dynamical under $\sigma 32$ -style control, with synthesis rising as free chaperone falls; the identity holds at median relative error 2.5×10^{-11} to 2.9×10^{-11} , and the $\sigma_0 \rightarrow 0$ limit reproduces the two-state field exactly. That limit is not itself a test of the identity: with synthesis off, dc/dt vanishes identically, $\nabla C = 0$, and both sides of the identity are exactly zero, so the reduction is checked against the two-state field rather than against the determinant. In the second system aggregate is split into reactive and sequestered compartments: median 1.5×10^{-12} , maximum 4.7×10^{-11} .

The corrected conditions hold at every fold state of both. Over the regulated system $|c_1| \geq 9.61 \times 10^{-3}$ at the standard controller and ≥ 0.132 at the stronger one; over

the sequestered system $|c_1| \geq 2.13 \times 10^{-2}$ across 46 fold states solved from 63 seeded attempts. In all three the remaining spectrum stays off the imaginary axis, the closest approach being $|\operatorname{Re} \lambda| = 1.62 \times 10^{-2}$. The $k_{seq} = 0$ and $\sigma_0 = 0$ settings are excluded from these margins rather than reported as passing cases: both make a row of the Jacobian identically zero, so $\det J = 0$ everywhere and the states a fold solve returns for them are not folds.

4.2 Growth dilution

Adding cell division leaves the theorem intact and removes the bound that made the capacity argument work.

Dilution does not affect binding, so the diluted field is $du/dt - \mu u$, $da/dt - \mu a$. H1 and H2 both survive, with total removal

$$R_{tot} = R + \mu(u, a) \cdot (u + a)$$

so $\det J = -(\nabla R_{tot} \times \nabla G_{dil})$ for any dilution law: verified at median relative error 1.2×10^{-10} for constant μ and 4.7×10^{-10} for burden-dependent μ , with the $\mu \rightarrow 0$ reduction exact. For constant μ the diluted Jacobian is $J - \mu I$, so the fold condition says exactly that μ is an eigenvalue of the undiluted Jacobian.

The removal ceiling does not survive. R_{tot} contains $\mu(u + a)$, unbounded in burden, so the bound $j_{crit} \leq C_{enz}$ fails once cells divide. At sufficiently large constant dilution the fold is destroyed outright: continuing in μ at the base parameters gives j_{crit} rising $0.1542 \rightarrow 0.2456$ as μ goes $0 \rightarrow 0.08$ while the fold state a^* diverges $0.265 \rightarrow 0.750$, and by $\mu = 0.10$ no fold exists. Past that point the low-burden branch no longer terminates and burden rises smoothly with influx. Across 25 parameter draws having a fold at $\mu = 0$, 23 lose it under constant dilution, at thresholds spanning 3.3 decades ($p10/p50/p90 = 0.0033/0.086/0.328$).

Burden-dependent growth restores it. With $\mu(u, a) = \mu_0/(1 + (u + a)/k_\mu)$ a fold exists at every rate tested up to $\mu_0 = 0.3$. The physiological content is that burden-dependent growth reduction preserves the fold in a regime where constant dilution destroys it, not that division alone is incompatible with a fold.

The coupling itself is well established. Expressing protein a cell does not need reduces its growth rate, and the reduction scales with expression level (Dekel and Alon 2005; Shachrai et al. 2010). Coarse proteome partitioning makes the mechanism explicit: raising any sector's share comes at the expense of the ribosomal sector, and the resulting growth law predicts the burden of heterologous expression quantitatively (Scott et al. 2010, 2014; Scott and Hwa 2023). Calibrating against a measured dosage-response — a 3.2% growth rate reduction when misfolded protein represents less than 0.1% of total cellular protein, measured in yeast (Geiler-Samerotte et al. 2011) and used here as a scaling assumption rather than a bacterial calibration — gives slope 32 per unit misfolded proteome fraction and linear arrest at fraction 0.03125, both bounds rather than point values. Under this calibration a fold survives in 30 of 30 networks with φ_{enz} confined to 0.072–0.147.

Four properties must be distinguished and are not interchangeable: existence of a local fold; boundedness of total burden; genuine runaway; and crossing a viability threshold. The results above concern the first two.

Corollary 3 (decomposition under division). *Splitting the critical influx into enzymatic and dilution parts,*

$$j_{crit} = \frac{C_{enz} \varphi_{enz}}{1 - \delta}, \quad \varphi_{enz} = \frac{R_{enz}(u^*, a^*)}{C_{enz}}, \quad \delta = \frac{R_{dil}(u^*, a^*)}{j_{crit}}$$

with (u^*, a^*) the fold state of the diluted system at that μ , and both factors dimensionless and in $[0, 1)$. Both are defined from that state, so the identity is a rearrangement rather than a derivation; it holds to 1.6×10^{-16} , and its content is how the two factors behave under division and not the identity itself.

At the base parameters φ_{enz} stays within 0.125–0.134 as μ runs $0 \rightarrow 0.08$, a $\pm 4\%$ band, while δ runs $0 \rightarrow 0.39$ and carries essentially all the variation. Division does not change the enzymatic condition for the fold; it multiplies the tolerable influx by $1/(1 - \delta)$. Loss of the fold is $\delta \rightarrow 1$, and across the draws above the median δ at the threshold is 0.35: the fold is typically lost once division does about a third of the disposal work. Since growth rate is part of disposal, an experiment that lets growth rate float is not holding disposal capacity fixed. The two factors are drawn on one axis in Fig. 2 because the contrast is the content: φ_{enz} is flat to within 7.5% of its own mean across the sweep while δ traverses most of its range.

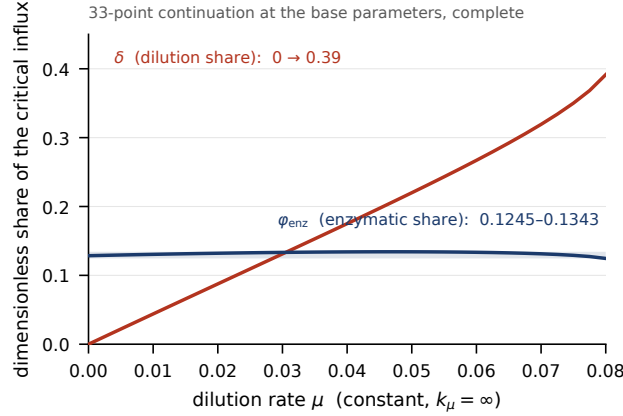


Fig. 2 Under division the enzymatic condition does not move; dilution does. A 33-point continuation at the base parameter set under constant dilution, plotted on one shared linear axis because the claim is a contrast between the two curves. Over μ from 0 to 0.080, φ_{enz} stays within 0.1245–0.1343, a full width of 7.5% of its own mean, while δ runs 0 to 0.3915 and carries essentially all the variation; j_{crit} rises 0.1542 to 0.2456 and a^* from 0.265 to 0.750. The band on φ_{enz} is a complete enumeration over the sweep, not an extremum over a subsample. The Corollary 3 identity holds to p99 1.74×10^{-16} and maximum 1.77×10^{-16} across the sweep.

4.3 Self-damaging machinery

In a cell, chaperones and proteases are translated by the error-prone apparatus, so raising j degrades the machinery that clears the damage.

We model this as $C_{enz}(load) = C_0/(1 + \varepsilon \cdot load)$ applied to both pools, with ε swept over four decades in two modes: $load = j$ and $load = u + a$. The second places capacity inside ∇R and ∇G themselves.

The identity holds in both modes at machine precision: gradient-normalised residual with floor 2.2×10^{-14} at $\varepsilon = 0$, and worst median 6.4×10^{-14} (influx mode) and 4.6×10^{-13} (burden mode) at $\varepsilon = 100$, where capacity has fallen to 16.7% and 1.8% of nominal. No correction to the algebraic form is required, for the reason given in Remark 1: the row operation needs (H1a) and H2 alone, and neither mentions da/dt .

What the coupling removes is the fixed nullcline. Under self-damage $\{G = 0\}$ moves with the load: (H1b) fails and Theorem 1.4 replaces Theorem 1.3, so the candidate solve grows from two equations in (u, a) to three in (u, a, j) . Corollary 1's saving applies to the frozen-capacity case.

It also removes the automatic transversality of Section 3.2, and puts a checkable scalar in its place. With $F_j = (1 - R_j - G_j, G_j)$ and $w = (1, 1 + \lambda)$,

$$w \cdot F_j = 1 - R_j + \lambda G_j,$$

so transversality is $R_j - \lambda G_j \neq 1$. This is the exact statement that transversality ceases to be free precisely when capacity feeds back on the influx.

Two quantities have to be kept apart. In the normalisation $w_1 = 1$ the left-hand side is a coefficient rather than a margin, since a left null vector rescales freely; the scale-invariant quantity is $|w \cdot F_j|/(\|w\| \cdot \|F_j\|)$. In burden mode the capacity factor carries no j , so R_j and G_j vanish identically and the coefficient is exactly 1 at all 46 states — that is the corollary of Section 3.2 reappearing, not a measurement, and those 46 values are not evidence. In influx mode the coefficient rises with the coupling rather than falling, from 1.000 at $\varepsilon = 0.01$ to 1.38 at $\varepsilon = 100$ over 53 states. The scale-invariant margin is the number to quote, and it is not uniformly comfortable: across the 99 states solved, from 200 attempts, its smallest value is 0.0070. Where it is small is predictable, and is the same geometry a third time. With $w = (1, 1 + \lambda)$ the cosine falls as $|\lambda|$ grows, and $|\lambda|$ grows as $\|\nabla G\|$ falls at fixed $\|\nabla R\|$ — which is the near-failure of (G1) again. Over these states log cosine correlates with log $\|\nabla G\|$ at +0.90 and with log $|\lambda|$ at -0.93, and the single smallest cosine and the single largest $|\lambda|$ are the same state. The small margins of Table 1, the two exceptions to the closed form, and the small cosines here are one geometric fact appearing three times, not three separate cautions.

Self-damage lowers the fold substantially without steepening the approach. Median j_{crit} relative to the frozen fold falls 0.999 / 0.990 / 0.925 / 0.640 / 0.322 across the influx ladder and 0.995 / 0.949 / 0.740 / 0.324 / 0.131 across the burden ladder. The critical-slowness exponent shows no detected shift: paired over the 19 networks carrying both values, median -0.4763 damaged against -0.4813 frozen, with a paired-difference distribution centred near zero (Wilcoxon $p = 0.312$).

Corollary 4 (square-root capacity ceiling). *In influx mode every removal flux carries a factor $1/(1 + \varepsilon j)$, so the necessary condition $j \leq C_0$ of the frozen model becomes $\varepsilon j^2 + j - C_0 \leq 0$, that is*

$$j \leq \frac{\sqrt{1 + 4\varepsilon C_0} - 1}{2\varepsilon} \longrightarrow \sqrt{C_0/\varepsilon} \quad \text{for } \varepsilon C_0 \gg 1.$$

A linear capacity ceiling becomes a square-root one: doubling the machinery buys only $\sqrt{2}$ in tolerable error rate once the machinery is itself error-prone. The bound is exact and is not binding in the range examined, with j_{crit}/j_{max} of median 0.039–0.186 and largest observed 0.623 over 8 draws, so it is a necessary condition and not the fold. Continuation with intermediate rungs recovers folds that a direct solve loses at large ε , with recovery counts that do not decrease monotonically in ε (influx 7/5/6/6/4, burden 6/7/7/4/2); a genuine loss of the fold would be monotone, so this is continuation failure rather than folds disappearing. Where a fold does solve it satisfies $\sin(\nabla R, \nabla G) < 2.0 \times 10^{-9}$ throughout that sweep. The consequence is that Corollary 4 is least checkable numerically in the regime where it would bind.

4.4 Robustness to the catalytic closure

Section 2.1 stipulates the catalytic step rather than deriving it. Nothing in Section 3 or in Sections 4.1 to 4.3 depends on that choice — the identity, the converse, the n -state form and the dilution corollaries follow from H1–H3, which hold for any closure preserving mass balance. The quantitative results of Sections 6 and 7 do depend on it, and this section reports by how much.

The alternative is the mechanistically standard one: catalytic flux proportional to the bound complex, so that $v_{ref} = k_{ref}C_f u_f / K_{CU}$ and likewise for the other three removal terms. It removes both stipulations at once. There is no second saturation, and cycles drawing on one pool now compete for it automatically, because the complexes they form appear in the same conservation law. The removal ceiling is unchanged at $C_{tot} + (\rho_U + \rho_A)D_{tot}$, since every complex concentration is bounded by its own pool total, so φ means the same thing under both and the comparison is well posed. Everything else is held fixed: the same draws, the same seeding from the same recorded states, one field of the parameter set changed.

Two things move, and by similar factors. Median φ at the fold rises from 0.0825 to 0.353, so the ceiling overestimates the tolerable influx twelvefold under the published rate laws and 2.8-fold under the alternative; paired over the 1991 networks yielding a fold state under both, the median ratio is 3.81 and 1516 of 1991 lie outside a factor of two. And the incidence of stability loss before the fold falls from 108 of 2766 traced networks to 18 of 2027. Fewer networks admit a solvable fold state at all under the alternative, 2027 against 2767 of the same 2884 draws, and that difference is reported rather than absorbed.

Neither movement is mysterious. The utilisation of nominal capacity at the fold is what φ aggregates, and removing a saturating factor from every removal flux raises it. The oscillatory region of Section 7 is characterised by κ_a , the Michaelis constant of aggregate clearance, which the alternative closure does not contain; what survives there is the weaker statement that the destabilising mechanism is saturation of clearance, whether that saturation enters through a Michaelis step or through competition for a finite pool.

The consequence for what this paper claims is stated where the claims are made. The ceiling overestimate is a property of the closure as much as of the model class, and the abstract, Section 6 and Section 10 say so. The identity and its corollaries are not.

5 Numerical verification

Two ensembles support the results. The **load grid** holds kinetics fixed at the base parameter set and sweeps nascent occupancy and rescue allocation, giving 325 folds; it describes how the fold moves within one network. The **kinetic box** is a Latin-hypercube ensemble of 5000 draws from a stipulated parameter box, of which 2884 admit a fold and 2767 yield a solvable fold state; it describes how the fold varies across networks.

Four denominators recur below and are not interchangeable, so they are fixed here once. **2884** is the draws admitting a fold, and is the population for anything read off the continuation itself. **2767** is those yielding a state that solves $\{G = 0, \det J = 0\}$, and is the population for anything evaluated at the fold. **2766** is the subset admitting a branch trace, and **2765** the subset whose orientation is evaluable. Each table and each count below names its own, and a figure quoted against one is not comparable with a figure quoted against another. Every quantity is recomputed from tracked generators and asserted by the accompanying test suite.

Table 2 Numerical verification of the identity and the solver. Each row names its own ensemble because the two are not interchangeable, and each maximum is reported beside a p99 because a maximum grows with the size of the population it is drawn from while a p99 does not. Row 1 compares two routes to one quantity: $\det J$ from the analytic Jacobian, whose partials are hand-derived, against $-(\nabla R \times \nabla G)$ from central differences of the two flux fields. The differentiation is independent; the closure solve is common to both, so the row tests the Jacobian and not the closure.

quantity	ensemble	median	p99	max
identity residual, gradient-normalised	load grid, 325	2.34×10^{-10}	9.67×10^{-10}	1.29×10^{-9}
$ G $ at fold states	load grid, 325	4.95×10^{-14}	9.40×10^{-10}	1.63×10^{-9}
direct solver against continuation	load grid, 325	3.03×10^{-7}	7.20×10^{-7}	7.56×10^{-7}
φ rebuilt from first principles	kinetic box, 2884	1.3×10^{-13}	2.98×10^{-9}	7.25×10^{-9}

Maxima over ensembles of different size are not directly comparable, and neither is comparable to a maximum computed at a third size; the p99 column is given for that reason, and at $n = 325$ it rests on a small number of upper-tail observations.

The identity residual sits at the central-difference floor. The parallelism residual is not zero at recorded states and should not be: those are bracketed approximations whose leading eigenvalue is of order 10^{-4} rather than 0. The correlation between log parallelism error and log leading eigenvalue is $+0.9960$ over the load grid, showing the residual to be bracket tolerance rather than failure of the identity, and the tightest-bracketed state in the complete population has $|\lambda| = 8.10 \times 10^{-10}$ with $\sin(\text{angle}) = 7.75 \times 10^{-9}$ (Fig. S1).

A residual between two expressions of one computation measures nothing, so the routes are worth naming. Row 1 of Table 2 differentiates by two independent routes over a closure solve common to both, and the four-product form of G_u in Section 3.3 is checked the same way, agreeing with a central-difference gradient to 2.6×10^{-10} across

the ensemble. Two checks in this paper are not of that kind, and both are stated where they occur: the $\sigma_0 \rightarrow 0$ reduction of Section 4.1, where both sides of the identity vanish identically and the comparison is made against the two-state field instead, and the self-damage check below.

One caution attaches to the self-damage check of Section 4.3. Because $du/dt = j - R - G$ holds pointwise and the central-difference operator is linear, that check reproduces the row operation exactly at any step size and carries no truncation term; measured slopes in the step size are -0.97 and -0.94 with no minimum over four decades, which is roundoff alone. It certifies that the implementation preserves mass balance. The analytic argument carries the result.

6 Where the fold sits

Writing $R = c_f s_{ref} + \rho_U d_f s_u + \rho_A d_f s_a$ against ceiling $= c_{tot} + \rho_U d_{tot} + \rho_A d_{tot}$, and $\varphi = j_{crit}/\text{ceiling}$, the saturation fractions at the fold answer how far below the ceiling it lies (Table 3). The decomposition is evaluated at re-solved fold states satisfying $\{G = 0, \det J = 0\}$ rather than where continuation stopped; of the 2884 draws admitting a fold, 2767 yield such a state and the remaining 117 are excluded and carry no entry. Evaluated instead at the recorded continuation states, over the same 2767 networks, φ reads 0.0803 against 0.0825, and the ceiling factor is twelvefold either way: the evaluation point moves the median by 2.7% and does not move the reported factor. Taken over all 2884 recorded states, including the 117 that do not re-solve, the median is 0.0769 and the factor thirteenfold — that difference is the excluded tail and not the evaluation point, and the two must not be quoted against each other.

Table 3 Where the fold sits relative to the capacity ceiling, at re-solved fold states over the 2767 networks that admit one. The saturation fractions are Michaelis factors, so a value of 1 would mean the machinery is running at maximum velocity.

at the fold, kinetic box, 2767 solved states	median
φ	0.0825
refolding saturation s_{ref}	0.180
soluble degradation saturation s_u	0.159
aggregate clearance saturation s_a	0.049
shortfall recovered by removing saturation	35.9%
shortfall recovered by removing sequestration	12.6%

At the fold the machinery delivers 8.3% of the flux its pools could support, its catalytic steps running at Michaelis factors of 5 to 18%, so the capacity ceiling is a correct bound that overestimates the tolerable influx about twelvefold under the stipulated catalytic closure — and 2.8-fold under the alternative of Section 4.4, which is why the factor is quoted with its closure wherever it is claimed. That clearance saturation eventually fails to balance aggregation is the qualitative content of the limited-capacity condition identified by Cotton et al. (2026);

the decomposition above quantifies how far short of saturation the failure occurs, and splits it.

Two questions must not be merged. Saturation dominates the *magnitude* of φ ; the *existence* of a turning point requires the superlinear aggregation that drives the free pools down at high burden. The counterfactuals answer the first only.

Dispersion is large. A nested design crossing 10 kinetic draws with a 7×7 load grid gives a between-to-within variance ratio for φ of 5.9, with per-network spread as large as $13.6 \times$ observed when both load coordinates sweep, against an $8.86 \times$ between-network spread — so φ is network-characteristic but not load-invariant. The saturation box was chosen deliberately wide, so part of the between-network spread is a property of the sampling. Adding $\sigma 32$ -style regulation widens the p5–p95 width of s_u from 0.890 to 0.968 over the 30 networks of that experiment rather than narrowing it, with regulated median s_u 0.323 against 0.169 unregulated; only 14 of 30 regulated networks converged against 24 unregulated, so that comparison is indicative. Over the kinetic box the p5–p95 width of s_u is 0.867.

Fig. 3 shows why the dispersion, rather than the median, is what a reader should take from this section: the p5–p95 spans cover most of the unit interval, so a single measured saturation fraction discriminates weakly. Its inset gives the reason no screening floor is applied to the low- s_a tail — the distribution runs smoothly across five decades with no gap, and the median slides continuously from 0.076 to 0.398 with whatever floor is imposed, so any floor is a free parameter moving a load-bearing number fivefold.

The claim this section supports is that the capacity bound overestimates by roughly an order of magnitude and that the fold is a computable constrained critical point — not that the fold occurs at any particular fraction of V_{max} .

7 Instabilities preceding the fold

The fold is not always the first loss of stability of the low-burden branch. Tracking $\text{tr } J$ along that branch from small j up to j_{crit} , the load grid shows no crossing in any of its 325 networks, with $\max(\text{tr } J)$ of median -0.338 and largest -0.243 . Over the kinetic box, of the 2767 solved fold states 2766 admit a branch trace, and 108 of those — 3.90% — reach $\text{tr } J = 0$ before the fold, with $\det J > 0$ at the crossing in all of them (minimum 1.5×10^{-6}), so the eigenvalues there are $\pm i\sqrt{\det J}$. In 104 of the 108 the branch’s influx maximum is the fold; in the remaining four a second singular point terminates the run the trace returns, so the crossing precedes a different branch’s endpoint and those are carried separately rather than counted. The first crossing occurs at a median of 0.83 of the way to j_{crit} , with deciles 0.40 / 0.61 / 0.83 / 0.95 / 0.995.

The crossing networks are not one population, and the split is the result. In 61 of the 108 the trace is still positive at the fold: stability is lost once and not regained, and the fold is not a stability boundary because the branch reaching it is already unstable. In the other 47 the trace returns below zero before the fold. Those networks carry an **instability window strictly interior to the stable branch**, opened by one crossing and closed by another, after which the branch is stable up to a fold that does terminate a stable state. Counting the zeros of $\text{tr } J$ confirms the parity: all 61 of

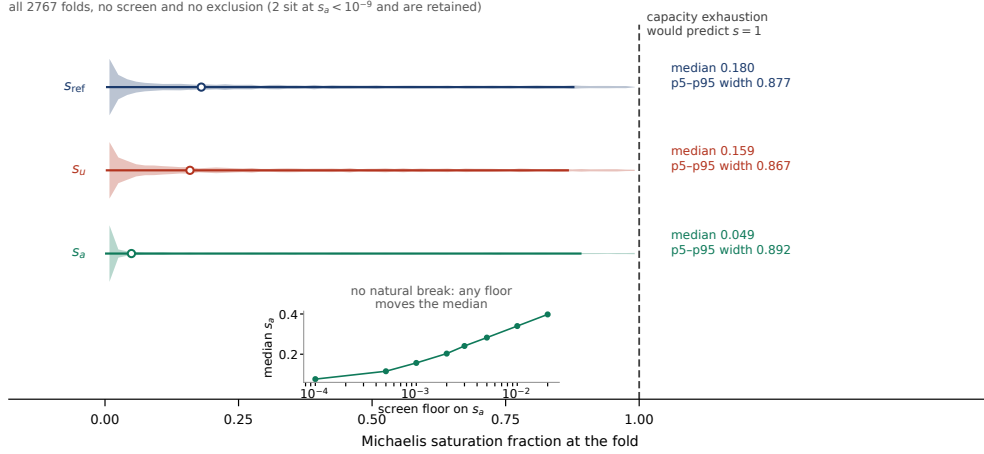


Fig. 3 Saturation of the clearance machinery at the fold, over the 2767 kinetic-box networks that yield a solvable fold state, with none excluded. The distributions are evaluated at the same re-solved states as the table in Section 6, so figure and table describe one population; the p5–p95 widths are 0.877 for s_{ref} , 0.867 for s_u and 0.892 for s_a . The point the figure carries is the width, not the centre: the spans are wide enough that a single measurement discriminates weakly against the capacity-exhaustion alternative. **Inset:** median s_a against an imposed lower screening floor, from 0.076 at a floor of 10^{-4} to 0.398 at 2×10^{-2} . The distribution runs smoothly across five decades with no gap, so any floor is a free parameter that moves a load-bearing median by a factor of five. No floor is applied anywhere in this paper.

the first group cross exactly once, and 38 of the 47 cross exactly twice. The remaining nine are branches the trace does not resolve — two are the multiplicity-ambiguous cases above, two leave the window through the high-burden side, and five stop short of low burden — and are reported rather than assigned. Over the 38 resolved windows the oscillatory band spans 0.0179 to 0.658 of j_{crit} , median 0.199, opening at a median 0.64 and closing at 0.95.

A crossing network and a non-crossing one are traced together in Fig. 4a: the first rises through zero at about 0.82 of the way to j_{crit} and stays positive to the fold, while the second approaches the fold from below without ever reaching it.

Integration confirms the instability independently of the branch computation, and confirms nothing beyond it. Perturbing each of the 104 crossing networks at its own $\text{tr } J$ maximum, all 104 amplify by more than tenfold and 102 reach the recording threshold, set at 10^{-2} of the state scale against a 10^{-6} perturbation — a ten-thousandfold amplification — with median $\max |\delta|/\delta_0$ of 9851; 205 non-crossing controls — every thirteenth of the 2658 non-crossing traces, taken systematically rather than at random — perturbed at their own $\text{tr } J$ maximum give 0 escapes and median 1.00, with no overlap between the distributions (Fig. 4b). All 309 states are exact equilibria. The integration also recovers the eigenvalue: growth exponent within 5% of $\max \text{Re}(\lambda)$ in 90 of 104, and the period of $\|\delta\|$ within 5% of π/ω in 47 of 49. What that establishes is that the crossing equilibria are unstable and the non-crossing controls are not, by a route sharing no failure mode with the branch computation. It does not establish runaway, and the word is avoided here for that reason: growth onto a stable limit cycle of finite amplitude reaches the same threshold as divergence, so the measurement records

departure from a small neighbourhood and cannot separate a bounded orbit from an unbounded trajectory. The rates by group are correspondingly indistinguishable — 60 of 61 against 42 of 43 — and that is a limit of the test rather than a result about the two groups. The reading applies equally to both: the terminal group’s 60 escapes are no more evidence of runaway than the window group’s 42.

What distinguishes them is the cubic term, and it is computable. At each of the 146 crossings the pair crosses the axis transversally, with $|d(\text{tr } J)/dj| \geq 0.293$; the sign is positive at every one of the 61 single crossings. Over the 38 windows whose two crossings are both resolved, the lower crossing in j is the one with index zero in all 38, and the sign is positive there and negative at the upper crossing in all 38, as the geometry requires. The nine unresolved window branches contribute one crossing each and are not a pair: seven of those nine carry a negative sign, which is what a trace that caught the closing crossing without its partner looks like, and they are reported here rather than folded into either count. The three groups are 61, 76 and 9, and sum to the 146. The first Lyapunov coefficient evaluates at 145 of the 146, with the sign stable across a factor of six in step size at every one of them. **120 of the 145 crossings are supercritical.** The instability windows are supercritical at both ends in 81 of their 85 crossings, so a stable limit cycle is born as the influx enters the window and is destroyed as it leaves: burden oscillates over a band of influx and then stops oscillating as the influx rises further, before the fold. Among the 61 single crossings the split is 39 supercritical against 21 subcritical, so about a third of that group loses stability hard, with no local attractor past the crossing.

Criticality at the two ends does not by itself say what happens between them, and that interior was integrated rather than interpolated. At seven influx values spanning each of the 38 resolved windows, the branch equilibrium was located and perturbed, and the full nonlinear system integrated in continuing blocks until the trajectory’s envelope stopped changing. Of 237 evaluable points 227 carry a settled oscillation, with period dispersion of median 8.2×10^{-3} over the final block; none diverges. **Cycles are present at every evaluable interior point in 34 of the 38 windows.** Amplitude varies smoothly along a window, spanning a median factor of 1.78 and at most 6.0 from end to end, with no jump between adjacent samples; the measured period exceeds the linear $2\pi/\omega$ by a median of 5.0%, which is the expected finite-amplitude correction. What this does not establish is that one connected branch of cycles spans the window: orbit continuation was not performed, and a smooth amplitude profile through seven samples is consistent with a single branch without demonstrating it. Of the remaining points, 29 at 11 networks had not settled inside the block budget and are reported rather than classified — one network, **draw4627**, has no evaluable point at all — and 10 at three networks settle instead onto a distinct fixed point with aggregate 10^{-5} to 10^{-3} of the unstable equilibrium’s. Those ten sit at the opening half of their windows and never the closing half, and subcriticality does not account for them: two of the three networks are supercritical at both crossings.

The crossings occupy a coherent region of parameter space (Fig. 4c). Relative to non-crossing networks, aggregate clearance is faster (ρ_A $7.5\times$ higher) and saturates more sharply (κ_a $11.7\times$ lower), and aggregate formation is growth-dominated rather than nucleation-dominated (α_g/α_n shifting from 3.7 to 33.6, with α_g alone only $1.67\times$

higher). Stated as a condition on the field rather than on the parameters, that description is the inequality $G_{af} \geq 0$ of Section 3.3, which holds at 97 of these 108 networks against 25.3% of the rest: the statement that aggregate is self-amplifying at fixed free machinery both destabilises the branch here and forbids the homeostatic degeneracy there. That configuration is a recognised destabilising one. Strong saturation of a degradation step destabilises a fixed point while weakening it restores stability, which has been identified as the determinant of oscillation in NF- κ B signalling, where the step is I κ B degradation (Krishna et al. 2006), and Michaelis-Menten degradation kinetics promotes oscillation across positive-feedback, negative-feedback and combined architectures (Xu and Qu 2012). In enzymatic oscillators driven by allosteric product activation the Michaelian character of product decay likewise governs period, amplitude and waveform (Goldbeter 2013). Here the autocatalytic partner is elongation, and no allosteric or transcriptional feedback is present.

This is a distinct mechanism from the oscillation question in the regulated heat-shock response, which turns on transcriptional feedback and delay (El-Samad et al. 2005; Zheng et al. 2016; Siv  ry et al. 2016). The networks above carry no regulation.

Three codimension-two degeneracies bound the region, and the ensemble approaches each of them at a different network, so they are three boundaries and not one organising centre. A fold with $\text{tr } J = 0$ is a Bogdanov–Takens point, and since $\det J = 0$ holds at every fold, the closest approach is the (G2) entry of Table 1: $|\text{tr } J| = 8.38 \times 10^{-4}$ over all folds, and 1.06×10^{-3} restricted to the networks unstable at their fold. That is also exactly where Theorem 2’s converse fails, since (G2) is the condition that breaks. A window closes when its two crossings collide, which is a tangency of $\text{tr } J$ to zero along the branch; the ensemble comes within $\max(\text{tr } J) = 4.56 \times 10^{-4}$ of it. And a crossing changes criticality where the Lyapunov coefficient vanishes, a Bautin point; the smallest $|l_1|$ over all crossings is 5.65×10^{-5} , at a window’s closing crossing. Closing crossings sit systematically nearer that boundary than opening ones, median $|l_1|$ 0.051 against 0.858.

The incidence figure is a property of the stipulated parameter box and is not offered as a prediction. The load grid’s cleanliness establishes that the base kinetic parameter set lies outside the oscillatory region, not that the region is small: the load grid holds kinetics fixed, and the crossing networks are distinguished by kinetic parameters.

8 Aggregation asymmetry in dividing *E. coli*

Unstressed *E. coli* accumulate aggregates at the old pole and lose reproductive capacity relative to new-pole progeny (Stewart et al. 2005; Lindner et al. 2008). Two features of the model class bear on this.

8.1 What the model class cannot do

No variant of the model examined places a stable aggregate-bearing attractor at a low growth-rate cost. Under constant dilution the system is bistable, with a 12.6- to 32.3-fold aggregate ratio between branches, but that regime predicts a growth-rate loss of exactly zero because μ cannot respond to burden. Under the calibrated hyperbolic law the system is monostable in four of six settings tested. Under linear arrest there is no

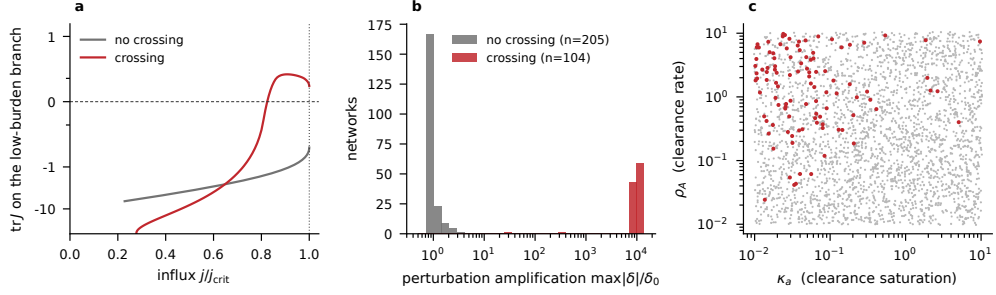


Fig. 4 The fold is not always the first loss of stability. **(a)** $\text{tr } J$ along the low-burden branch against j/j_{crit} , on a symmetric-log axis, for one crossing network and one non-crossing network. Both exemplars are chosen by rule rather than by inspection: the crossing case is the network nearest the median crossing position among those whose branch is resolved below $0.3 j_{crit}$, and the non-crossing case is the network nearest the population median of $\max(\text{tr } J)$. The crossing exemplar happens to be one of the 61 that do not recover, so its trace stays positive to the fold; a window network would return below zero before reaching it. **(b)** Amplification of a perturbation applied at each network’s own $\text{tr } J$ maximum, integrating the full nonlinear system. Of the 108 crossing networks the 104 whose branch maximum is the fold were integrated: all grow by more than tenfold and 102 leave the linear neighbourhood, at a median amplification of 9851; the 205 non-crossing controls give 0 escapes and a median of 1.00. The distributions do not overlap, and this panel shares no failure mode with the branch computation in (a). It does not, however, separate a stable limit cycle from divergence — escape is a ten-thousandfold amplification and both produce it — which is what the Lyapunov classification in the text is for. **(c)** Where the crossing networks sit in the sampled (κ_a, ρ_A) plane. Relative to the rest they have ρ_A $7.5\times$ higher and κ_a $11.7\times$ lower: clearance that is fast and sharply saturating.

bounded high-burden state: μ reaches exactly zero beyond k_μ , switching dilution off and removing the bound. Adding a sequestered compartment — aggregate drained into an inert pool that neither nucleates nor occupies chaperone or protease — produces 12 bistable cells against a control’s 3, but none of 384 qualified cells predicts a loss within the measured range, every bistable cell predicting 0.482, 0.508, 0.954 or 1.000. Inverting the growth law, the high state is at least 7.5 to 254 times more aggregate-laden than the measured cell, and the upper figure is a lower bound because the 1.000 values are the linear-arrest law clamping at states 3.4 to 43 times past arrest. A real old-pole cell carries a visible inclusion body and divides at 96–99% of normal.

8.2 The observation does not require bistability

The old-pole cell does not occupy a second basin. It inherits a physical inclusion body at each division, which is a continuously renewed perturbation rather than an attractor, and a monostable system under a renewed perturbation has a stationary offset with no separatrix.

The two-state diluted model under the calibrated hyperbolic law, with asymmetric partitioning of aggregate at division and no additional state variable, reproduces the measured lineage difference. Of 728 parameter settings, 709 settle, and 43 score within the measured range with the partitioning active. The control is exact: all 66 settings at symmetric partitioning give a lineage difference of 0.0 with standard deviation 0.0, which follows algebraically, since symmetric partitioning into half the volume leaves concentration unchanged and is what the $-\mu x$ term already encodes.

8.3 A requirement on the aggregate load

The scored quantity is Lindner et al.’s $\Delta(GR_{old} - GR_{new})/GR_{mean}$, whose aggregate-attributable part is 1.04–1.78%, obtained from the reported lineage difference of $-3.95 \pm 0.5\%$ and the reported aggregate-attributable share of 30–40%.

Partitioning is set by one physical quantity. The visible focus is indivisible and passes entirely to one daughter, but the model’s a is total aggregate, of which the focus is a share β . The old-pole daughter then receives $(1 + \beta)a$ and the new-pole daughter $(1 - \beta)a$ after division into half the volume, which is the scalar rule at $f_{eff} = (1 + \beta)/2$. With $k_\mu = 0.03125/p_{qc}$, the scored quantity is $32 \times (B_{old} - B_{new})$ as a proteome fraction exactly, so the quality-control proteome share and the growth rate enter only through the stationary load. Inverting gives a requirement scaling as $1/\beta$:

Table 4 The old-pole aggregate load required by the account of Section 8.3, as a function of the share β of aggregate held in the visible polar focus, over the same range Fig. 5 plots. The final column divides the only aggregate load that has been measured, 5–10% of total protein in $\Delta rpoH$ at 30 °C, by the requirement, so an entry is the factor by which the measured load exceeds it; the columns pair the extremes, so they give the widest and narrowest separations consistent with both intervals.

β	f_{eff}	required old-pole aggregate (% of proteome)	$\Delta rpoH$ load / required load
1.00	1.000	0.047 – 0.080	62× – 214×
0.75	0.875	0.061 – 0.104	48× – 165×
0.50	0.750	0.091 – 0.157	32× – 110×
0.25	0.625	0.182 – 0.314	16× – 55×
0.05	0.525	0.913 – 1.570	3.2× – 11×

The damping factor for the daughter’s relaxation during its own generation is 0.346–0.355, with weak β dependence.

Table 4 spans the same β range the figure plots. At $\beta = 1.00$, where the whole aggregate is in the focus, the requirement sits 62× to 214× below the $\Delta rpoH$ load; at $\beta = 0.05$ it is 3.2× to 11× below it. The requirement is therefore smaller than the only load ever measured, and the informative comparison is how far below it falls.

No source we could find bounds β under unstressed growth. Winkler et al. (2010) report aggregate mass and focus size separately, and under heat stress, which is neither the condition at issue nor a focus-share measurement. The requirement is therefore a β -indexed family (Fig. 5), and the direction is the informative part: the less of the aggregate held in the focus, the more aggregate the account requires, and the closer the requirement moves to the only load anyone has measured. Tomoyasu et al. (2001) report 5–10% of total protein aggregated in $\Delta rpoH$ mutants at 30 °C and no detectable aggregation in $rpoH^+$ cells at the same temperature; wild type is a bound rather than a value, and the reported mutant figures are not a limit of detection for the assay.

Two measurements close the question: the wild-type aggregate fraction under unstressed growth, and the share of that aggregate held in the polar focus, by quantitative fluorescence of the focus against total aggregated protein by sedimentation in the same cells.

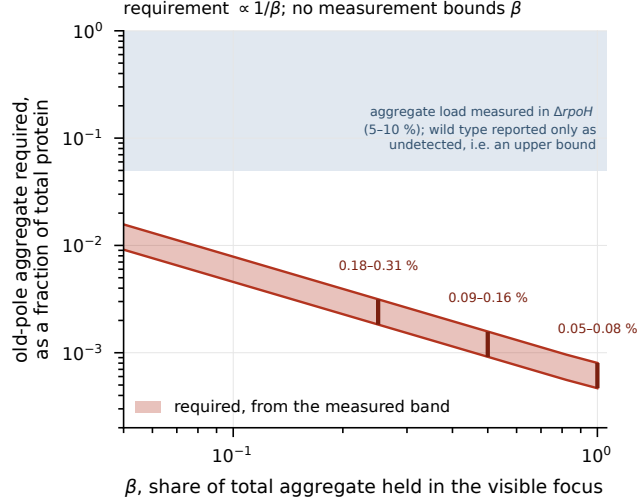


Fig. 5 The old-pole aggregate load required by the account of Section 8.3, as a function of the focus share β , over the same range as the table there. The requirement scales as $1/\beta$: at $\beta = 1.00$ it is 0.0467–0.0803% of the proteome and at $\beta = 0.05$ it is 0.9126–1.5696%. The interval at each β comes from the aggregate-attributable part of the measured lineage difference, 1.04–1.78%. The shaded band is the only aggregate load anyone has measured, 5–10% of total protein in $\Delta\rho H$ at 30 °C; no measurement bounds β , so no vertical line is drawn.

9 Predictions

Mathematical. The identity $\det J = -(\nabla R \times \nabla G)$, its n -state form, the converse under (G1)–(G3), the $\mu \rightarrow 0$ and $\sigma_0 \rightarrow 0$ reductions, the $J - \mu I$ form under constant dilution, the (φ_{enz}, δ) decomposition, and Corollary 4 follow from H1–H3 and no observation bears on them.

The fold sits far below saturation. The alternative is capacity exhaustion, in which the fold occurs as $s \rightarrow 1$. The observable is a ratio, so testing it does not require absolute copies-per-cell calibration. Two cautions: the dispersion in s_u is wide (p5–p95 width 0.867 over the kinetic box), so a single measurement discriminates weakly, and adding regulation widens rather than narrows it.

The fold moves under nascent load. Since nascent chains consume chaperone capacity without contributing influx, raising the load of perfectly folding protein should lower the tolerable damage influx; a model handling the two loads independently predicts no shift. Direction was correct in 67 of 68 settings over a 100-fold ladder, with median shift $1.22\times$. Growth rate sets the load coordinate and contributes about $1.6\times$ on its own against the $\sim 1.8\times$ contrast expected, so the experiment requires externally fixed growth rate: chemostat or turbidostat, not batch culture.

Oscillation before collapse, and then not. In networks with fast, sharply saturating aggregate clearance against growth-dominated aggregation, oscillations in aggregate burden should appear as the damage influx rises. The predicted period of the burden itself is $2\pi/\omega$ with $\omega = \sqrt{\det J}$ at the crossing; the perturbation norm peaks at twice that rate, every π/ω , which is the quantity Section 5’s verifier recovers and not the one an experiment would measure.

Where both crossings are supercritical the oscillation is sustained rather than transient, carried by a stable limit cycle found at every evaluable interior influx in 34 of the 38 windows integrated, and the sharper prediction is what happens next: the cycle is destroyed at the second crossing, so burden oscillates over a band of influx and then stops, with the branch stable again up to a fold that does terminate it. Non-monotone loss and recovery of stability under monotone loading is not what a clearance model is expected to do, and the band is wide enough to look for — a median 20% of the influx range up to j_{crit} , opening near 0.64 and closing near 0.95. Where the single crossing is subcritical, about a third of that group, there is no cycle and the loss of stability is hard.

Oscillation costs less mean burden than its amplitude suggests. Over the 227 settled orbits the time-averaged aggregate is a median 1.15 times the equilibrium value it oscillates about, p95 3.04 and at most 4.40, while the peak-to-peak swing is a median 2.4 times that equilibrium. Mean and amplitude rise together — Spearman 0.69, and the upper 5% of each is the same four networks — so these are one quantity seen twice rather than two independent observables, and an experiment resolving large swings should expect a mean raised modestly rather than not at all.

The old-pole aggregate load. Section 8.3 turns the measured lineage growth difference into a requirement on the aggregate a dividing *E. coli* old-pole cell must carry, indexed by the share β held in the visible focus. It is the one prediction here that a single pair of measurements can refute, and the requirement sits $62\times$ to $214\times$ below the only aggregate load anyone has measured, so the two quantities needed are the wild-type aggregate fraction under unstressed growth and the focal share in the same cells.

A ruler, not a test. Approaching the fold, relaxation time scales as $\tau \sim (j_{crit} - j)^{(-1/2)}$ (fitted slope 0.5077, $r^2 = 1.0000$). This is generic to folds, so confirming it selects this model over nothing in particular. Its use is as an instrument for locating the fold in an experiment that then asks whether the fold moves.

The measurement the theory most wants. Whether growth arrest under misfolding burden is complete or asymptotic. Under the hyperbolic form, dilution keeps bounding the high-burden state and loss of proteostasis is a reversible switch; under linear arrest, μ reaches exactly zero and the high state is a runaway. The available dosage-response constrains the slope only below 0.1% misfolded protein, about 1.5 decades below the arrest burden where the two forms diverge, so it cannot adjudicate. That single functional form decides whether losing proteostasis in a dividing cell is recoverable.

10 Discussion

The fold of a mass-balanced clearance network is not something one has to find by sweeping. Given the removal flux and the aggregate nullcline, it is where their gradients become parallel, and the critical influx is the removal flux evaluated there. The statement requires no fitted parameter, holds for any number of states, survives growth dilution, and survives capacity that degrades with the load it carries. Two

things bound what it delivers: the candidate set can have more than one element, and the fold is not always the first loss of stability.

What it does not do is predict a number without measured parameters. The dispersion in φ across the kinetic ensemble is large enough, and wide enough after calibration and after adding regulation, that the quantitative predictions above are ratio comparisons rather than point predictions. The twelvefold overestimate is also a property of the stipulated catalytic closure and not of the model class: the mechanistically standard alternative gives 2.8-fold, and Section 4.4 reports what else moves with it. The identity and its corollaries do not depend on that choice. This is a structural result rather than a predictive one.

Santra et al. (2019) reach a tipping point in a proteostasis model by chaperone titration under cumulative damage, which is the same qualitative claim about how collapse arises. The present contribution is the exact flux-geometric condition for where it occurs, and the classification of what else can happen first.

The Hopf result identifies a regime worth looking for. Oscillatory dynamics of aggregate burden in a clearance network requires no regulatory circuit — only fast, sharply saturating clearance against autocatalytic growth, which is the classical saturable-degradation destabilisation acting on an aggregation system. Whether real quality-control networks occupy that region is open, and Section 9 states the signature.

The prediction a reader can act on is Section 8.3. If the aggregate burden of an aging *E. coli* old-pole cell falls outside the interval required at the measured focal share, the account of the aging asymmetry given here is wrong. Neither of the two measurements needed has been attempted as far as we can determine.

Statements and Declarations

Competing interests. The author declares no competing financial or non-financial interests.

Funding. No funds, grants or other support were received for conducting this study.

Ethics approval, consent, and human or animal participants. Not applicable. The study is theoretical and computational and involved no human participants, no animals, and no newly generated experimental data.

Author contributions. Sole author: conceptualisation, derivation, software, analysis, and writing.

Data and code availability. All analysis code, parameter configurations, and the test suite asserting every numerical quantity reported here are at <https://github.com/khatvangi/proteostasis-law-theory> under the MIT licence, archived at <https://doi.org/10.5281/zenodo.21794565> (concept DOI, resolving to the latest version).

One limitation of the deposit is stated rather than glossed. The two ensembles — the load grid and the kinetic box — were generated by a run whose working tree was uncommitted at launch, and no patch was recorded, so the launch-time source is

not recoverable. The deposit therefore archives the tracked source of each generating experiment, with per-file checksums, in place of a commit identifier alone; for the structural-check experiment the recorded commit contained no code at all, and its snapshot is taken from the earliest commit that does, which the index marks. The gap lies between that source and the raw ensembles. Every quantity in this paper is recomputed from those raw outputs by the tracked code above, under the test suite, so no reported number depends on the unrecorded state.

References

- Antoneli F, Golubitsky M, Jin J, Stewart I (2025) Homeostasis in input-output networks: structure, classification and applications. *Math Biosci* 384:109435. doi:10.1016/j.mbs.2025.109435
- Araujo RP, Liotta LA (2018) The topological requirements for robust perfect adaptation in networks of any size. *Nat Commun* 9:1757.
- Brennan GS, Thompson TB, Oliveri H, Rognes ME, Goriely A (2023) The role of clearance in neurodegenerative diseases. *SIAM J Appl Math* S172–S198. doi:10.1137/22M1487801
- Brualdi RA, Ryser HJ (1991) *Combinatorial matrix theory*. Cambridge University Press, Cambridge.
- Cohen SIA, Linse S, Luheshi LM, Hellstrand E, White DA, Rajah L, Otzen DE, Vendruscolo M, Dobson CM, Knowles TPJ (2013) Proliferation of amyloid- β 42 aggregates occurs through a secondary nucleation mechanism. *Proc Natl Acad Sci USA* 110:9758–9763.
- Cotton MW, Goriely A, Klenerman D, Meisl G (2026) A universal phase-plane model for in vivo protein aggregation. *J Chem Phys* 164(7):075101. doi:10.1063/5.0312752
- Dekel E, Alon U (2005) Optimality and evolutionary tuning of the expression level of a protein. *Nature* 436:588–592.
- El-Samad H, Kurata H, Doyle JC, Gross CA, Khammash M (2005) Surviving heat shock: control strategies for robustness and performance. *Proc Natl Acad Sci USA* 102:2736–2741.
- Ferrell JE (2016) Perfect and near-perfect adaptation in cell signaling. *Cell Syst* 2:62–67.
- Ferrone F (1999) Analysis of protein aggregation kinetics. *Methods Enzymol* 309:256–274.
- Geiler-Samerotte KA, Dion MF, Budnik BA, Wang SM, Hartl DL, Drummond DA (2011) Misfolded proteins impose a dosage-dependent fitness cost and trigger a cytosolic unfolded protein response in yeast. *Proc Natl Acad Sci USA* 108:680–685.
- Goldbeter A (2013) Oscillatory enzyme reactions and Michaelis-Menten kinetics. *FEBS Lett* 587:2778–2784. doi:10.1016/j.febslet.2013.07.031
- Golubitsky M, Stewart I (2017) Homeostasis, singularities and networks. *J Math Biol* 74:387–407.
- Golubitsky M, Wang Y (2020) Infinitesimal homeostasis in three-node input-output networks. *J Math Biol* 80:1163–1185.

- Hipp MS, Kasturi P, Hartl FU (2019) The proteostasis network and its decline in ageing. *Nat Rev Mol Cell Biol* 20:421–435.
- Huang Z, Golubitsky M (2022) Classification of infinitesimal homeostasis in four-node input-output networks. *J Math Biol* 84:25.
- Jin J, Rempala GA (2026) Infinitesimal homeostasis in mass-action systems. *J Math Biol* 92(3):35. doi:10.1007/s00285-026-02352-y
- Khammash M (2021) Perfect adaptation in biology. *Cell Syst* 12:509–521.
- Knowles TPJ, Waudby CA, Devlin GL, Cohen SIA, Aguzzi A, Vendruscolo M, Terentjev EM, Welland ME, Dobson CM (2009) An analytical solution to the kinetics of breakable filament assembly. *Science* 326:1533–1537.
- Krishna S, Jensen MH, Sneppen K (2006) Minimal model of spiky oscillations in NF- κ B signaling. *Proc Natl Acad Sci USA* 103:10840–10845. doi:10.1073/pnas.0604085103
- Lin X, Antoneli F, Wang Y (2026) Automated classification of homeostasis structure in input-output networks. *Bull Math Biol* 88(7):113. doi:10.1007/s11538-026-01679-3
- Lindner AB, Madden R, Demarez A, Stewart EJ, Taddei F (2008) Asymmetric segregation of protein aggregates is associated with cellular aging and rejuvenation. *Proc Natl Acad Sci USA* 105:3076–3081.
- Ma W, Trusina A, El-Samad H, Lim WA, Tang C (2009) Defining network topologies that can achieve biochemical adaptation. *Cell* 138:760–773.
- Madeira JLO, Antoneli F (2024) Homeostasis in networks with multiple inputs. *J Math Biol* 89:17.
- Meisl G, Michaels TCT, Linse S, Knowles TPJ (2017) Scaling behaviour and rate-determining steps in filamentous self-assembly. *Chem Sci* 8:7087–7097.
- Nijhout HF, Best J, Reed MC (2014) Escape from homeostasis. *Math Biosci* 257:104–110.
- Reed M, Best J, Golubitsky M, Stewart I, Nijhout HF (2017) Analysis of homeostatic mechanisms in biochemical networks. *Bull Math Biol* 79:2534–2557.
- Reed MC, Duncan W, Nijhout HF, Best J, Golubitsky M (2018) Homeostasis despite instability. *Math Biosci* 300:130–137.
- Santra M, Dill KA, de Graff AMR (2019) Proteostasis collapse is a driver of cell aging and death. *Proc Natl Acad Sci USA* 116:22173–22178.
- Scott M, Gunderson CW, Mateescu EM, Zhang Z, Hwa T (2010) Interdependence of cell growth and gene expression: origins and consequences. *Science* 330:1099–1102.
- Scott M, Klumpp S, Mateescu EM, Hwa T (2014) Emergence of robust growth laws from optimal regulation of ribosome synthesis. *Mol Syst Biol* 10:747.
- Scott M, Hwa T (2023) Shaping bacterial gene expression by physiological and proteome allocation constraints. *Nat Rev Microbiol* 21:327–342.
- Shachrai I, Zaslaver A, Alon U, Dekel E (2010) Cost of unneeded proteins in *E. coli* is reduced after several generations in exponential growth. *Mol Cell* 38:758–767.
- Sivéry A, Courtade E, Thommen Q (2016) A minimal titration model of the mammalian dynamical heat shock response. *Phys Biol* 13:066008.
- Stewart EJ, Madden R, Paul G, Taddei F (2005) Aging and death in an organism that reproduces by morphologically symmetric division. *PLoS Biol* 3:e45.

Thompson TB, Meisl G, Knowles TPJ, Goriely A (2021) The role of clearance mechanisms in the kinetics of pathological protein aggregation involved in neurodegenerative diseases. *J Chem Phys* 154:125101.

Tomoyasu T, Mogk A, Langen H, Goloubinoff P, Bukau B (2001) Genetic dissection of the roles of chaperones and proteases in protein folding and degradation in the *Escherichia coli* cytosol. *Mol Microbiol* 40:397–413.

Wang Y, Huang Z, Antoneli F, Golubitsky M (2021) The structure of infinitesimal homeostasis in input-output networks. *J Math Biol* 82:62.

Winkler J, Seybert A, König L, Pruggnaller S, Haselmann U, Sourjik V, Weiss M, Frangakis AS, Mogk A, Bukau B (2010) Quantitative and spatio-temporal features of protein aggregation in *Escherichia coli* and consequences on protein quality control and cellular ageing. *EMBO J* 29:910–923.

Xu L, Qu Z (2012) Roles of protein ubiquitination and degradation kinetics in biological oscillations. *PLoS ONE* 7:e34616. doi:10.1371/journal.pone.0034616

Zheng X, Krakowiak J, Patel N, Beyzavi A, Ezike J, Khalil AS, Pincus D (2016) Dynamic control of Hsf1 during heat shock by a chaperone switch and phosphorylation. *eLife* 5:e18638.